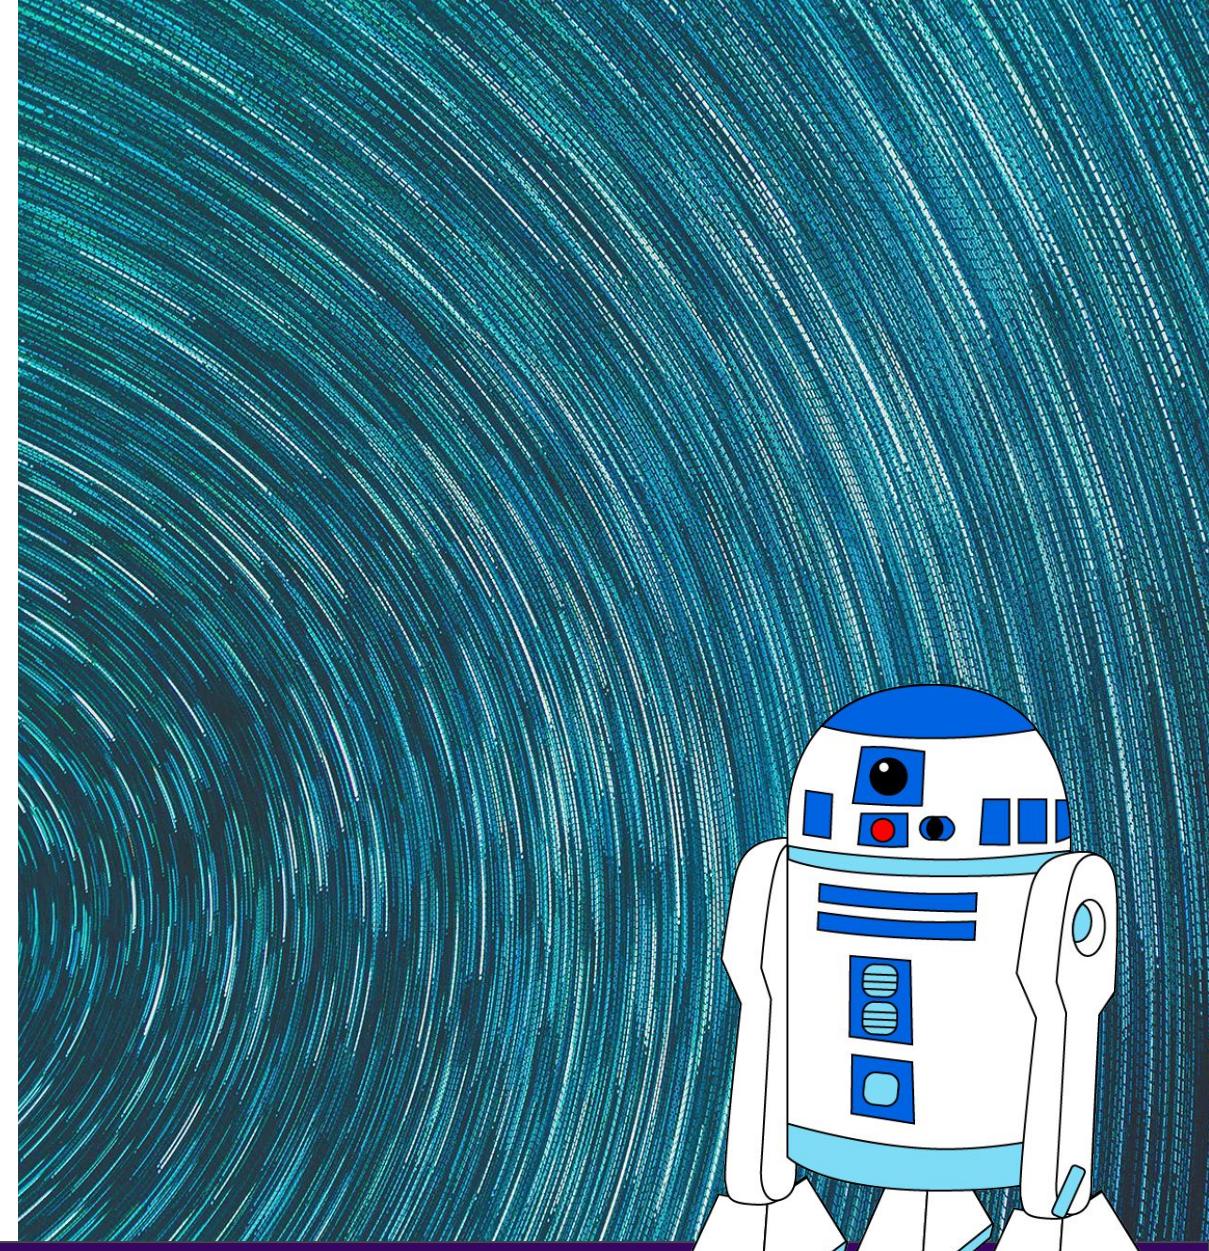


CIS 421/521:  
ARTIFICIAL INTELLIGENCE

# Games and Adversarial Search

Professor Chris Callison-Burch



# Games: Outline of Unit

- Part 1: Games as Search
  - Motivation
  - Game-playing AI successes
  - Game Trees
  - Evaluation Functions
- Part II: Adversarial Search
  - The Minimax Rule
  - Alpha-Beta Pruning

# May 11, 1997

may 11th game 6 : may 11 @ 3:00PM EDT | 19:00 GMT kasparov 2.5 deep blue 2.5

Home ▶ The match ▶ The players ▶ The technology ▶ Community

## Deep Blue Wins 3.5 to 2.5

### KASPAROV vs DEEP BLUE the rematch

With a dramatic victory in Game 6, Deep Blue won its six-game rematch with Champion Garry Kasparov ▶

OVERVIEW ▶ EVENT COVERAGE ▶ MATCH NEWS ▶ MAIN STORIES

 Commentary  
George Plimpton on chess, Kasparov, and the limitations of computers  
▶ Read the article

 Commentary  
Vishwanathan Anand on the legacy of Kasparov vs. Deep Blue  
▶ Read the article

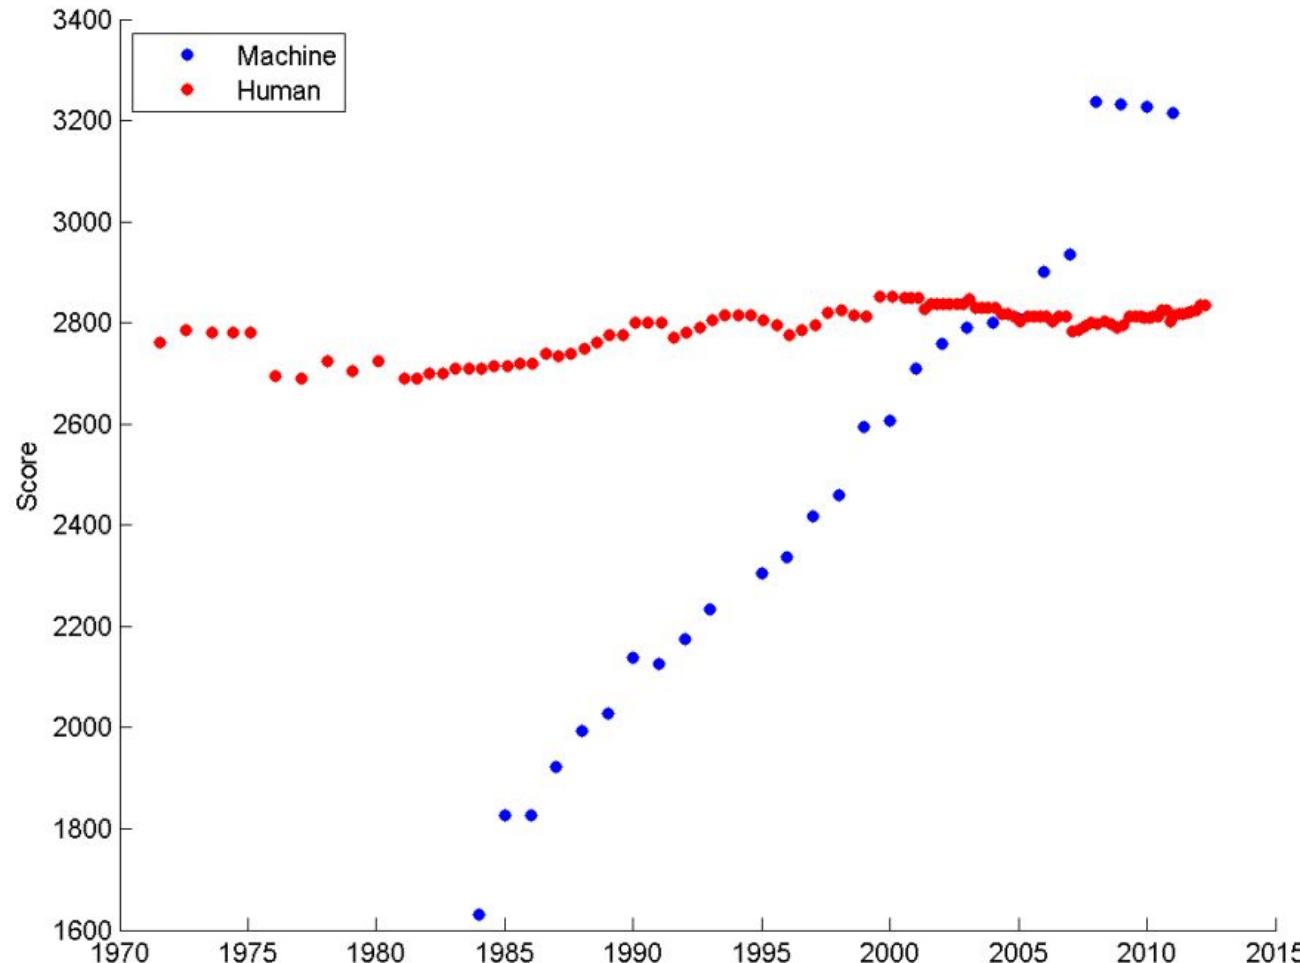
 Club Kasparov  
Visit the virtual home of the world's greatest chess player.

 Guest essays  
Thoughts on chess, computers, and what it all means  
▶ Read the essays...

 Community  
During the rematch, more than 20,000 people from 120 countries joined the community to talk about the match.

 Clips from the rematch  
Video footage from the games  
▶ Highlights from the games

# Ratings of human and computer chess champions



<https://srconstantin.wordpress.com/2017/01/28/performance-trends-in-ai/>

# AlphaGo seals 4-1 victory over Go grandmaster Lee Sedol

DeepMind's artificial intelligence astonishes fans to defeat human opponent and offers evidence computer software has mastered a major challenge

Steven Borowiec

Tuesday 15 March 2016 06.12 EDT



This article is 6 months old

Shares

613

Save for later



The world's top Go player, Lee Sedol, lost the final game of the Google DeepMind challenge match.  
Photograph: Yonhap/Reuters

Google DeepMind's AlphaGo program triumphed in its final game against South Korean Go grandmaster Lee Sedol to win the series 4-1, providing further evidence of the landmark achievement for an artificial intelligence program.

Lee started Tuesday's game strongly, taking advantage of an early mistake by AlphaGo. But in the end, Lee was unable to hold off a comeback by his opponent, which won a narrow victory.

# The Simplest Game Environment

- *Multiagent*
- *Static*: No change while an agent is deliberating.
- *Discrete*: A finite set of percepts and actions.
- *Fully observable*: An agent's sensors give it the complete state of the environment.
- *Strategic*: The next state is determined by the current state and the action executed by the agent and the actions of one other agent.

# Key properties of our games

1. Two players alternate moves
  2. Zero-sum: one player's loss is another's gain
  3. Clear set of legal moves
  4. Well-defined outcomes (e.g. win, lose, draw)
- o Examples:
    - Chess, Checkers, Go,
    - Mancala, Tic-Tac-Toe, Othello ...

# More complicated games

- Most card games (e.g. Hearts, Bridge, etc.) and Scrabble
  - Stochastic, not deterministic
  - Not fully observable: lacking in perfect information
- Real-time strategy games
  - Continuous rather than discrete
  - No pause between actions, don't take turns
- Cooperative games

# Pac-Man



<https://youtu.be/-CbyAk3Sn9I>

# Formalizing the Game setup

1. Two players: **MAX** and **MIN**; **MAX** moves first.
  2. **MAX** and **MIN** take turns until the game is over.
  3. Winner gets award, loser gets penalty.
- o Games as *search*:
    - *Initial state*: e.g. board configuration of chess
    - *Successor function*: list of (move,state) pairs specifying legal moves.
    - *Terminal test*: Is the game finished?
    - *Utility function*: Gives numerical value of terminal states.  
e.g. win ( $+\infty$ ), lose ( $-\infty$ ) and draw (0)
    - **MAX** uses search tree to determine next move.

# How to Play a Game by Searching

- **General Scheme**

1. Consider all legal successors to the current state ('board position')
2. Evaluate each successor board position
3. Pick the move which leads to the best board position.
4. After your opponent moves, repeat.

- **Design issues**

1. Representing the 'board'
2. Representing legal next boards
3. Evaluating positions
4. Looking ahead

# Hexapawn: A very simple Game

- Hexapawn is played on a 3x3 chessboard



- Only standard pawn moves:

1. A pawn moves forward one square onto an empty square
2. A pawn “captures” an opponent pawn by moving diagonally forward one square, if that square contains an opposing pawn. The opposing pawn is removed from the board.

# Hexapawn: A very simple Game

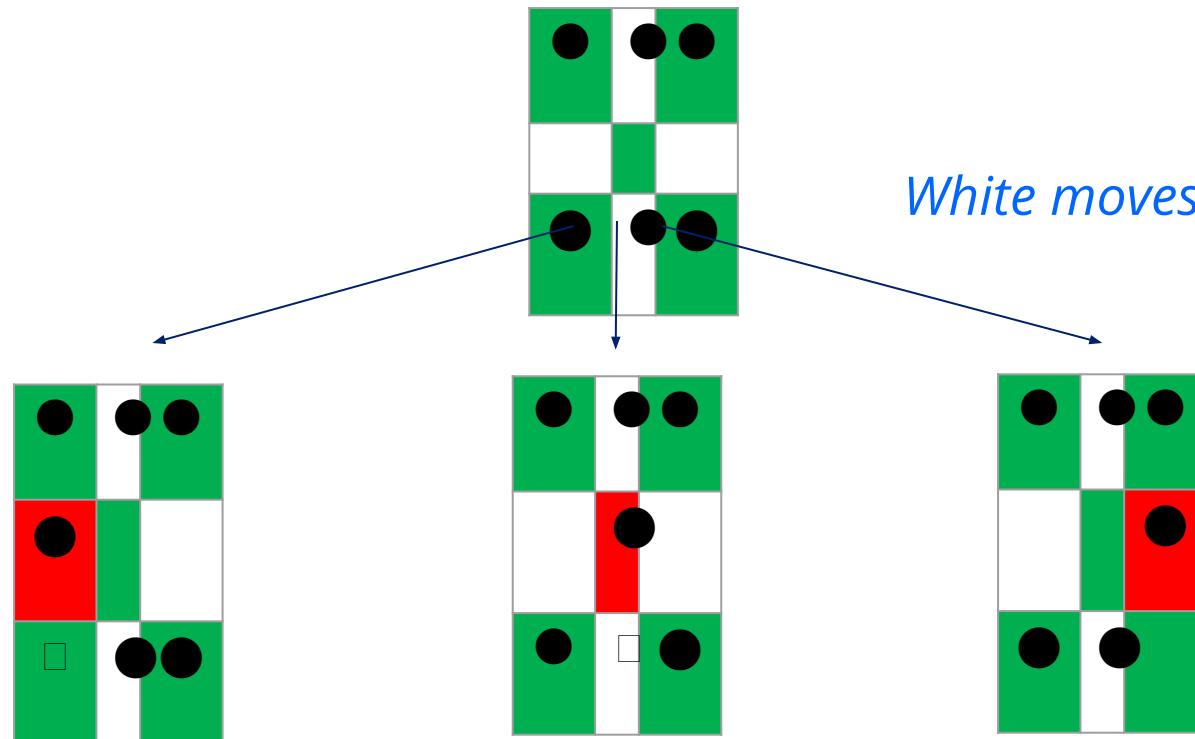
- Hexapawn is played on a 3x3 chessboard



- Player  $P_1$  wins the game against  $P_2$  when:
  - One of  $P_1$ 's pawns reaches the far side of the board, or
  - $P_2$  cannot move because no legal move is possible.
  - $P_2$  has no pawns left.

*(Invented by Martin Gardner in 1962, with learning “program” using match boxes.)*

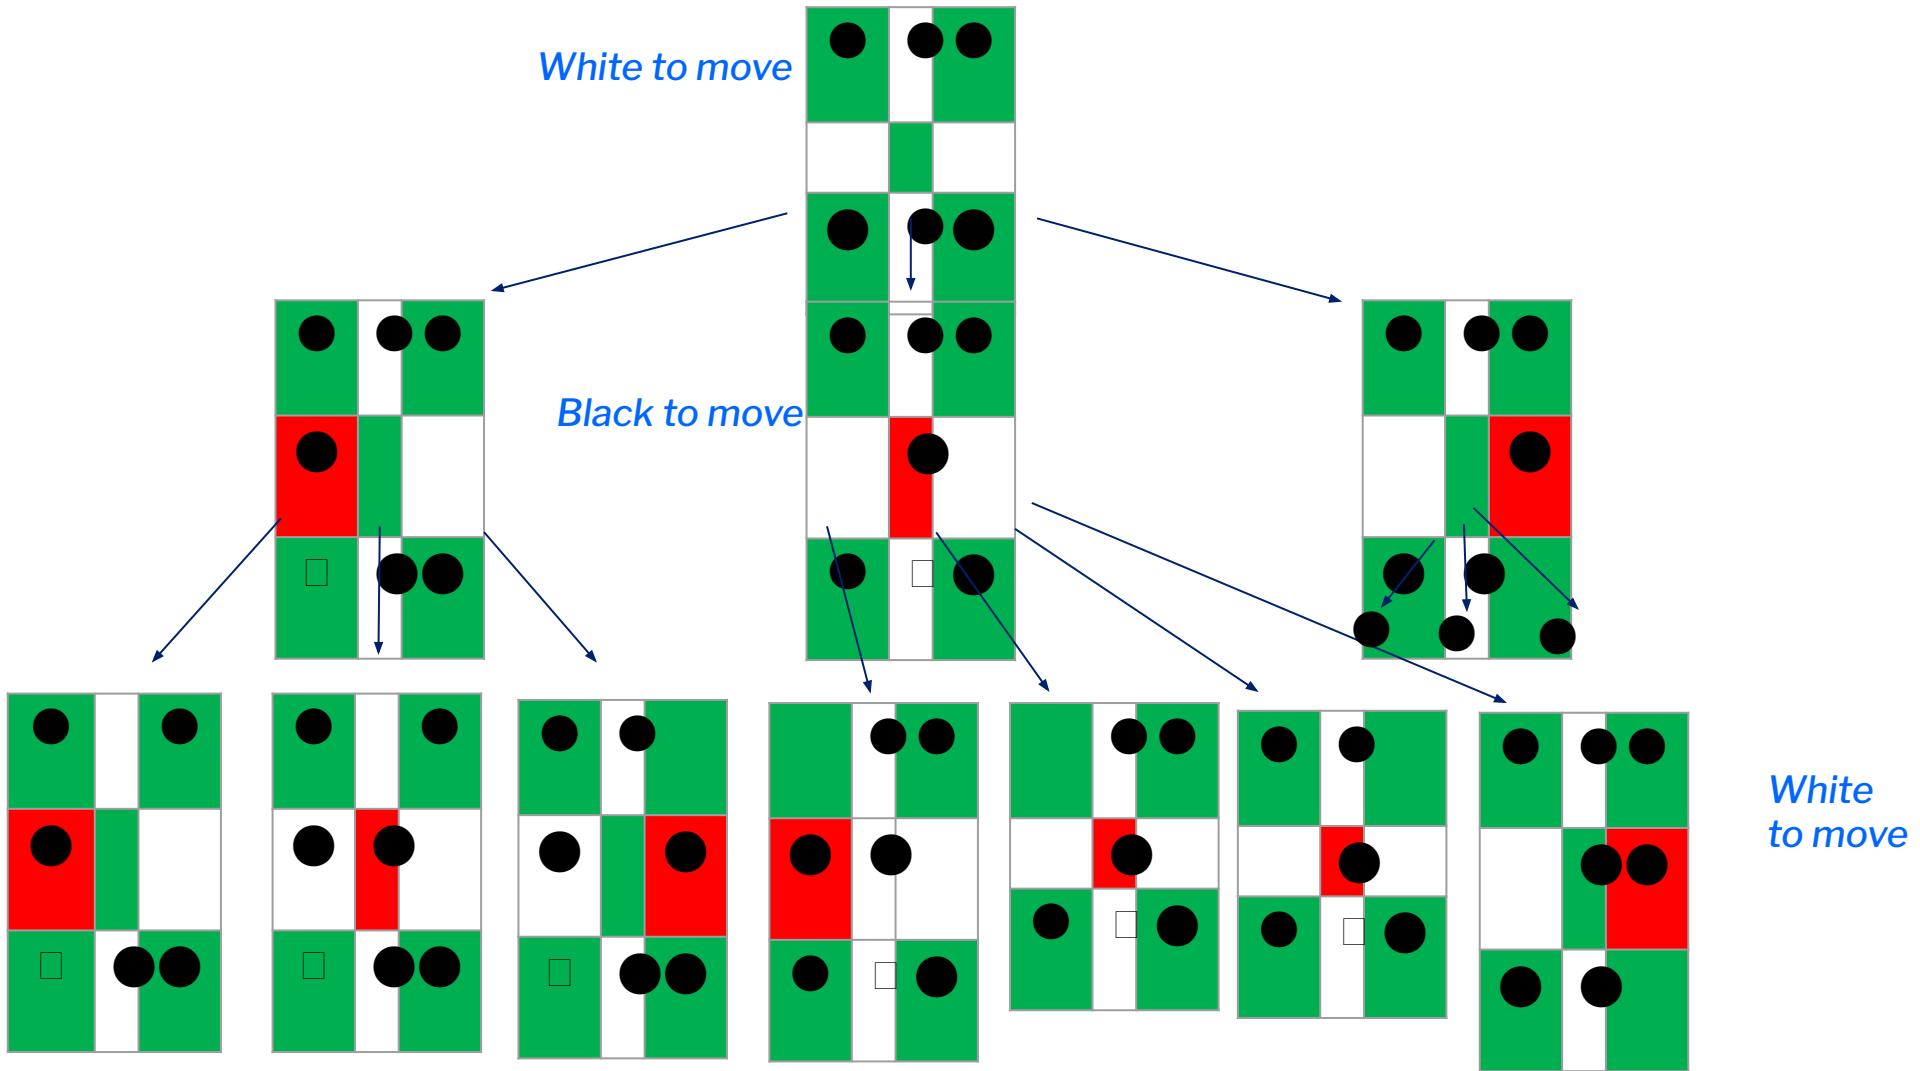
# Hexapawn: Three Possible First Moves



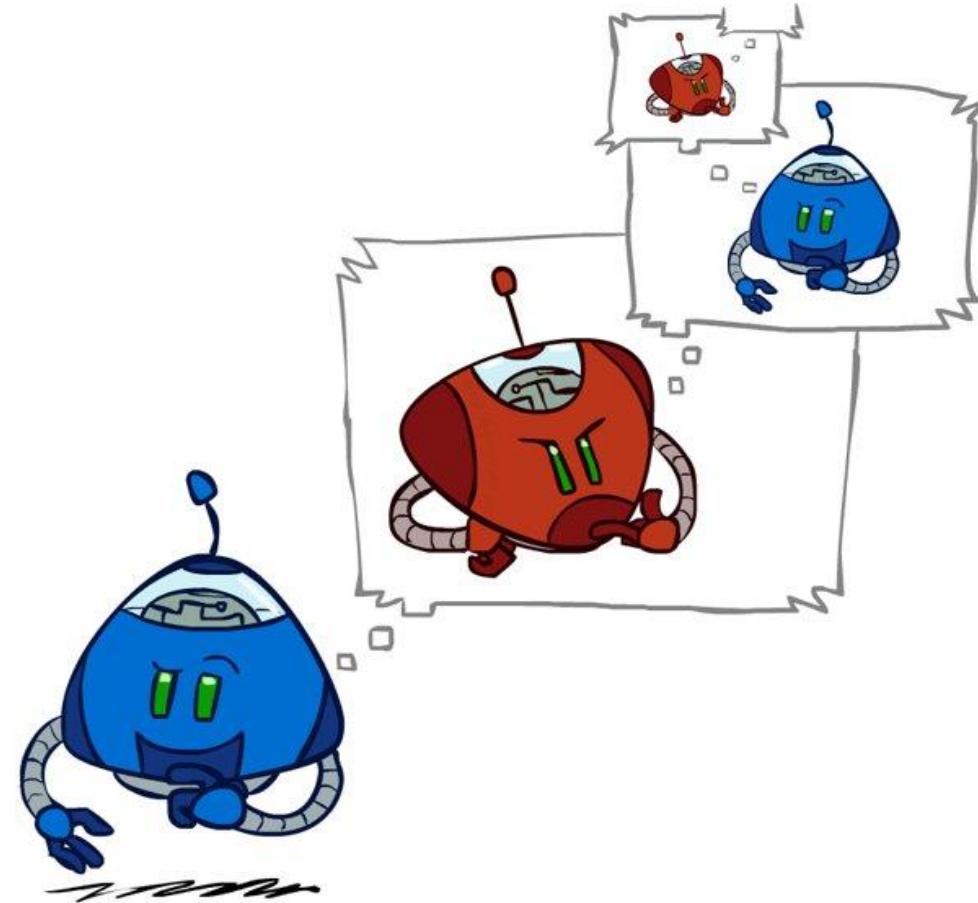
# Game Trees

- **Represent the game problem space by a tree:**
  - Nodes represent 'board positions'; edges represent legal moves.
  - Root node is the first position in which a decision must be made.

# Hexapawn: Simplified Game Tree for 2 Moves



# Adversarial Search



# Battle of Wits



FANDANGO  
MOVIECLIPS

<https://www.youtube.com/watch?v=rMz7JBRbmNo>

# MAX & MIN Nodes : An egocentric view

- Two players: MAX, MAX's opponent MIN
- *All play is computed from MAX's vantage point.*
- When MAX moves, MAX attempts to MAXimize MAX's outcome.
- When MAX's opponent moves, they attempt to MINimize MAX's outcome.
  - WE TYPICALLY ASSUME MAX MOVES FIRST:
- Label the root (level 0) MAX
- Alternate MAX/MIN labels at each successive tree level (*ply*).
- *Even levels* represent turns for MAX
- *Odd levels* represent turns for MIN

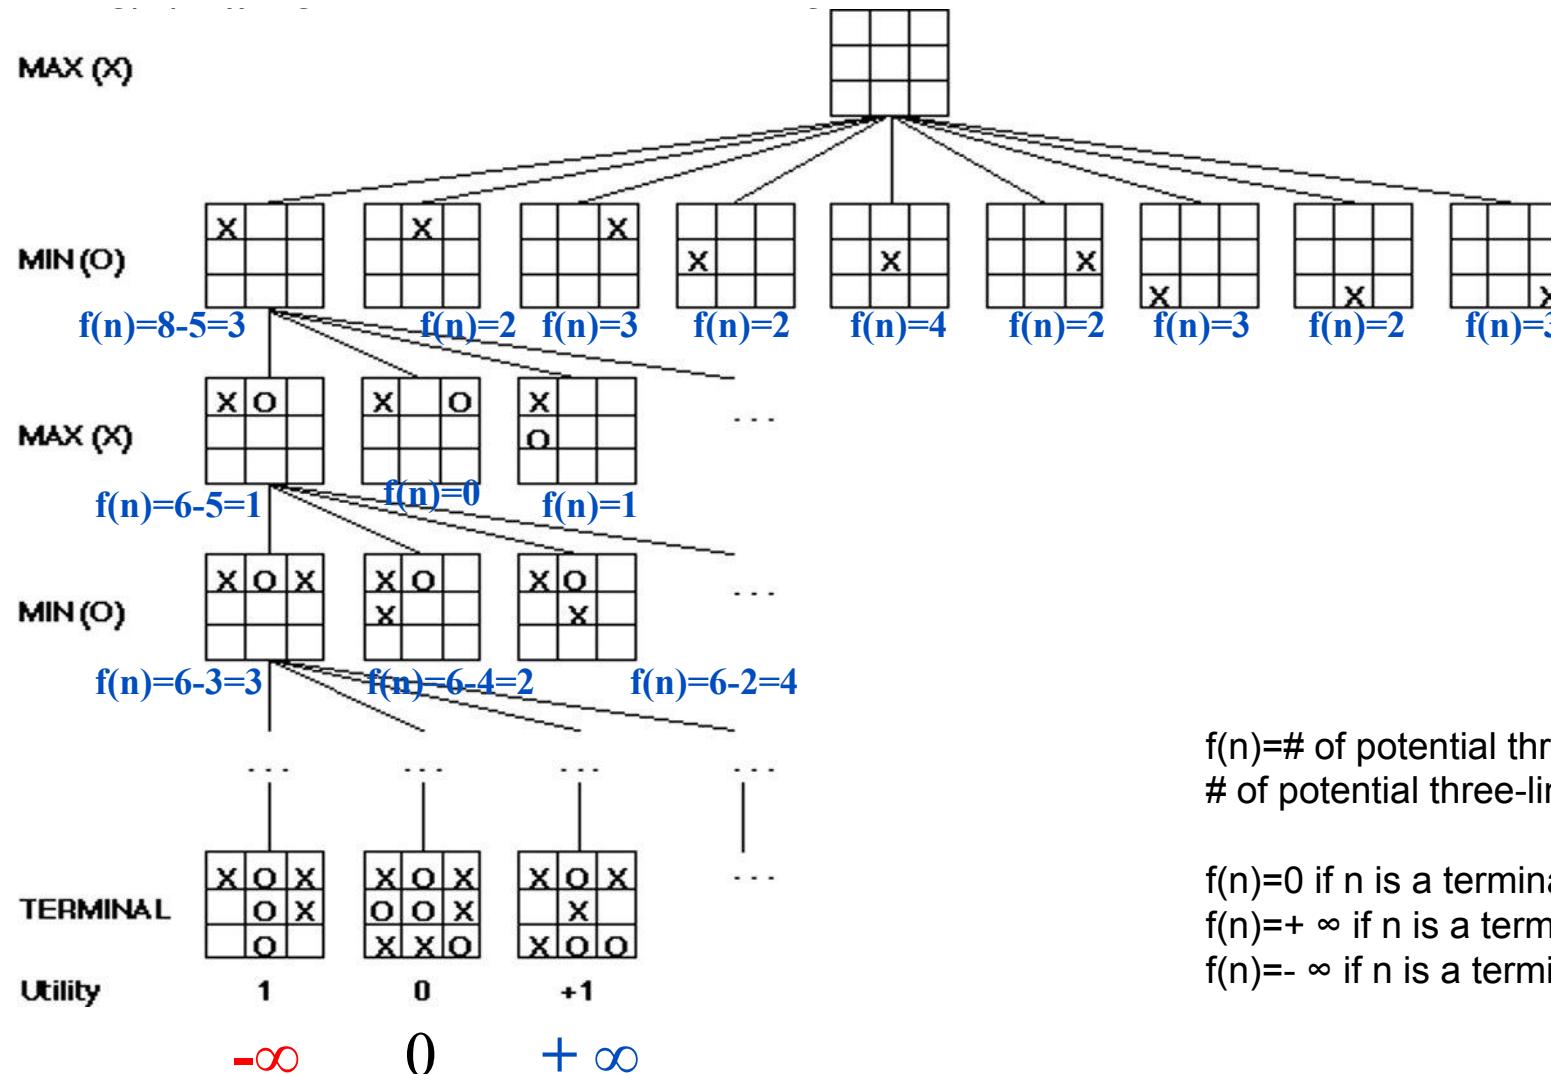
# Game Trees

- Represent the game problem space by a tree:
  - Nodes represent ‘board positions’; edges represent legal moves.
  - Root node is the first position in which a decision must be made.
- Evaluation function  $f$  assigns real-number scores to ‘board positions’ *without reference to path*
- Terminal nodes represent ways the game could end, labeled with the desirability of that ending (e.g. win/lose/draw or a numerical score)

# Evaluation functions: $f(n)$

- Evaluates how good a ‘board position’ is
- Based on *static features* of that board alone
- Zero-sum assumption lets us use one function to describe goodness for both players.
  - $f(n) > 0$  if MAX is winning in position  $n$
  - $f(n) = 0$  if position  $n$  is tied
  - $f(n) < 0$  if MIN is winning in position  $n$
- Build using expert knowledge,
  - Tic-tac-toe:  $f(n) = (\# \text{ of 3 lengths open for MAX}) - (\# \text{ open for MIN})$

# A Partial Game Tree for Tic-Tac-Toe



$f(n) = \# \text{ of potential three-lines for X} - \# \text{ of potential three-line for O}$

$f(n)=0$  if n is a terminal tie  
 $f(n)=+\infty$  if n is a terminal win  
 $f(n)=-\infty$  if n is a terminal loss

# Chess Evaluation Functions

- Claude Shannon argued for a chess evaluation function in a 1950 paper
- Alan Turing defined function in 1948:  
 $f(n) = (\text{sum of A's piece values}) - (\text{sum of B's piece values})$
- More complex: weighted sum of *positional* features:  
 $\sum w_i \text{feature}_i(n)$
- Deep Blue had >8000 features

Pawn	1.0
Knight	3.0
Bishop	3.25
Rook	5.0
Queen	9.0

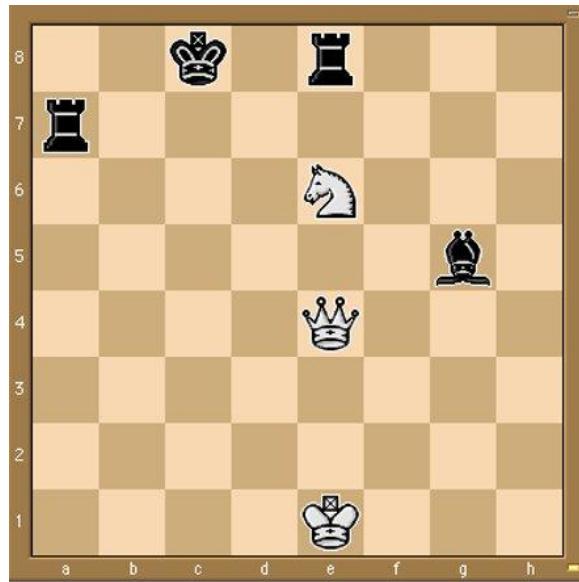
Type equation here.

Pieces values for a simple  
Turing-style evaluation function often  
taught  
to novice chess players

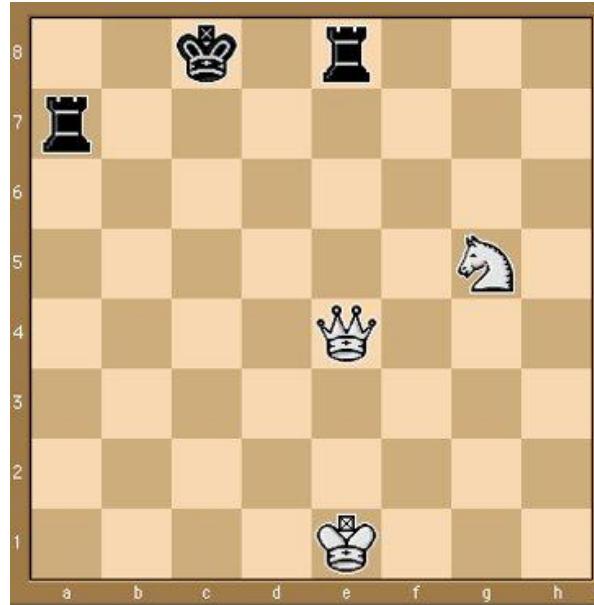
**Positive:** rooks on open files, knights in  
closed positions, control of the center,  
developed pieces

**Negative:** doubled pawns, wrong-colored  
bishops in closed positions, isolated pawns, pinned pieces  
*Examples of more complex features*

# Some Chess Positions and their Evaluations

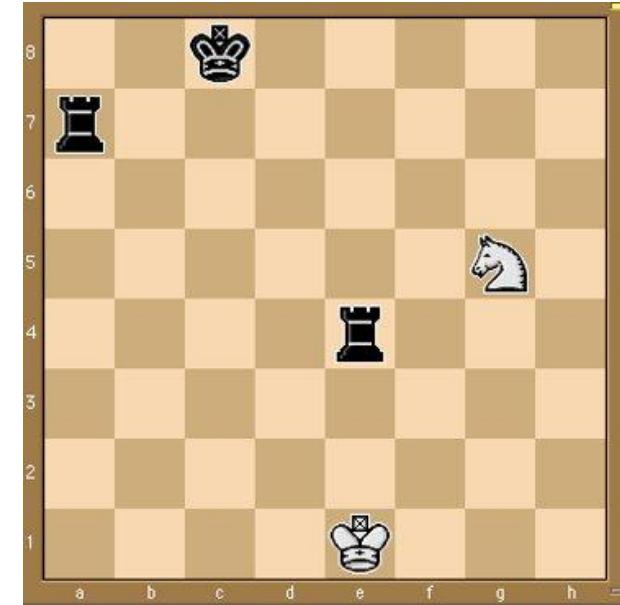


White to move  
 $f(n) = (9+3)-(5+5+3.25)$   
 $= -1.25$



... Nxe5??  
 $f(n) = (9+3)-(5+5)$   
 $= 2$

So, considering our opponent's possible responses would be wise.



*Uh-oh:* Rxg4+  
 $f(n) = (3)-(5+5)$   
 $= -7$

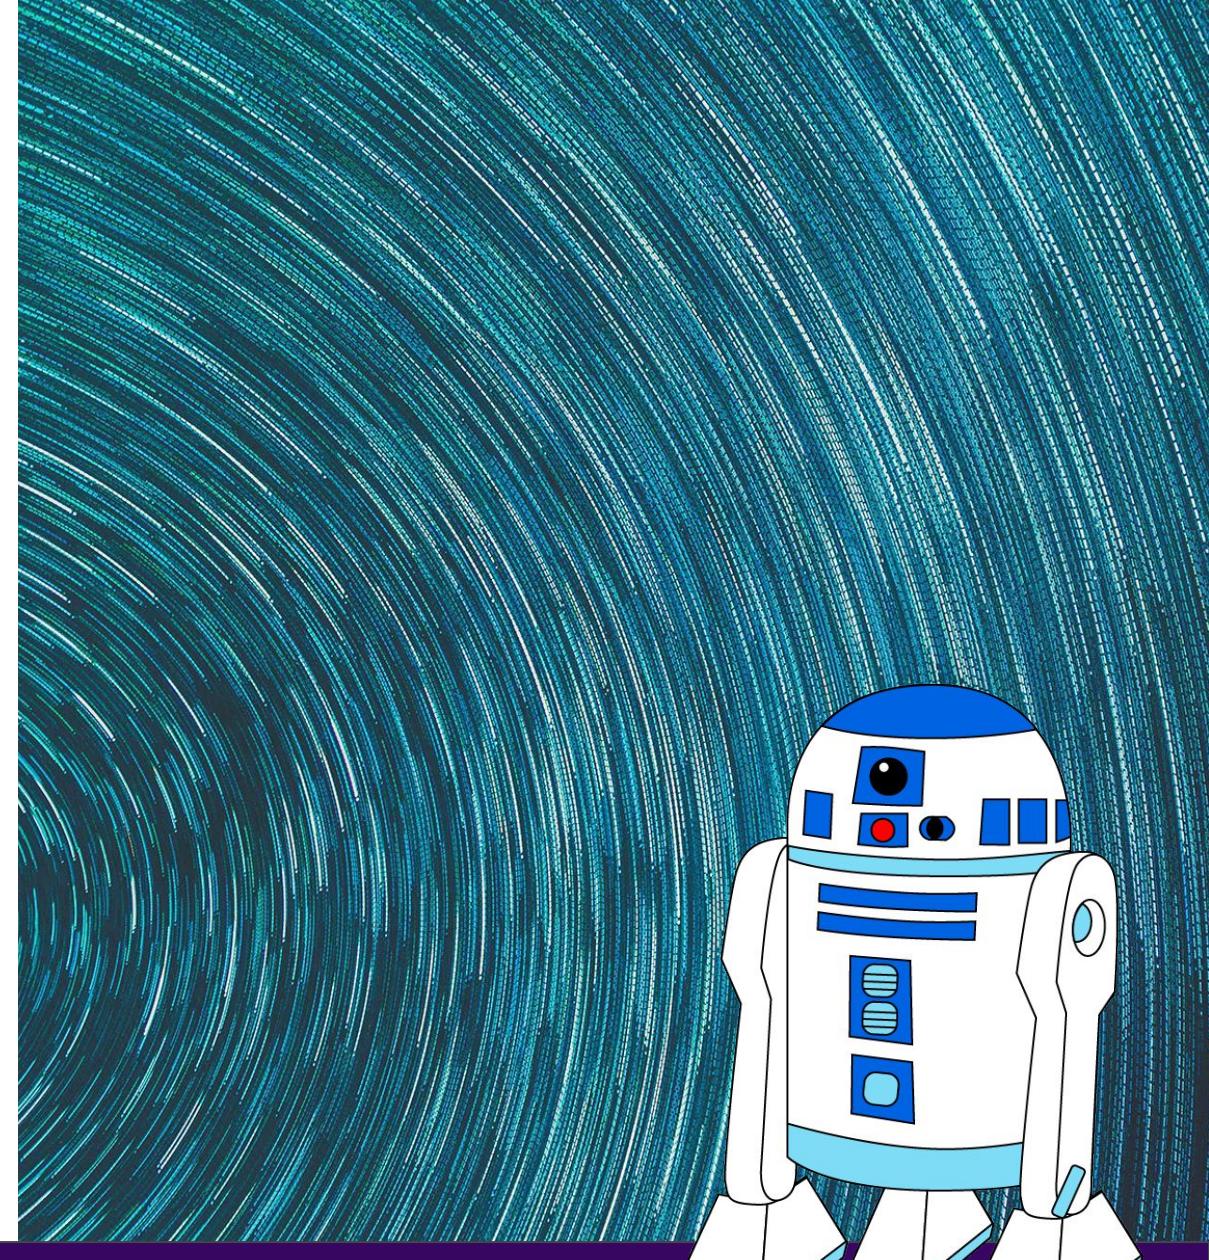
*And black may force checkmate*

# The Minimax Rule (AIMA 5.2)

CIS 421/521:  
ARTIFICIAL INTELLIGENCE

# Minimax and Alpha-Beta Pruning

Professor Chris Callison-Burch



# The Minimax Rule: “Don’t play hope chess”

- Idea: Make the best move for MAX *assuming that MIN always replies with the best move for MIN*
- Easily computed by a recursive process
  - The **backed-up value** of each node in the tree is determined by the values of its children:
    - For a **MAX** node, the backed-up value is the **maximum** of the values of its children (*i.e. the best for MAX*)
    - For a **MIN** node, the backed-up value is the **minimum** of the values of its children (*i.e. the best for MIN*)

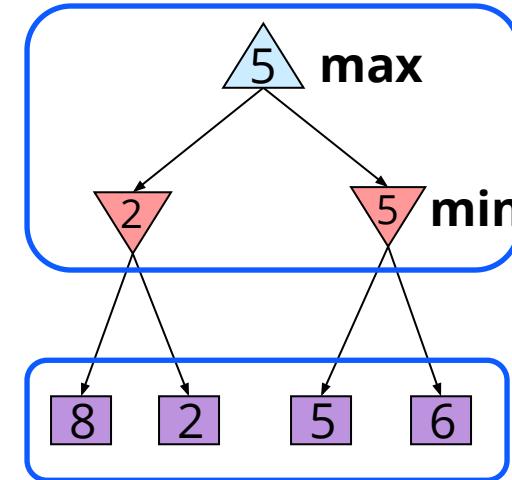
# The Minimax Procedure

- Until game is over:
  1. Start with the current position as a MAX node.
  2. Expand the game tree a fixed number of *ply*.
  3. Apply the evaluation function to the leaf positions.
  4. Calculate back-up values bottom-up.
  5. Pick the move assigned to MAX at the root
  6. Wait for MIN to respond

# Adversarial Search (Minimax)

- Minimax search:
  - A state-space search tree
  - Players alternate turns
  - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary

**Minimax values:  
computed recursively**



**Terminal values:  
part of the game**

# Minimax Implementation

```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v,
                 max-value(successor))
    return v
```



# Alpha-Beta Pruning

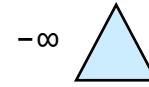
- During Minimax, keep track of two additional values:
  - $\alpha$ : MAX's current *lower* bound on MAX's outcome
  - $\beta$ : MIN's current *upper* bound on MIN's outcome
- MAX will never allow a move that could lead to a worse score (for MAX) than  $\alpha$
- MIN will never allow a move that could lead to a better score (for MAX) than  $\beta$
- Therefore, stop evaluating a branch whenever:
  - When evaluating a MAX node: a value  $v \geq \beta$  is backed-up
    - MIN will never select that MAX node
  - When evaluating a MIN node: a value  $v \leq \alpha$  is found
    - MAX will never select that MIN node

For  $\alpha$  think "at least"

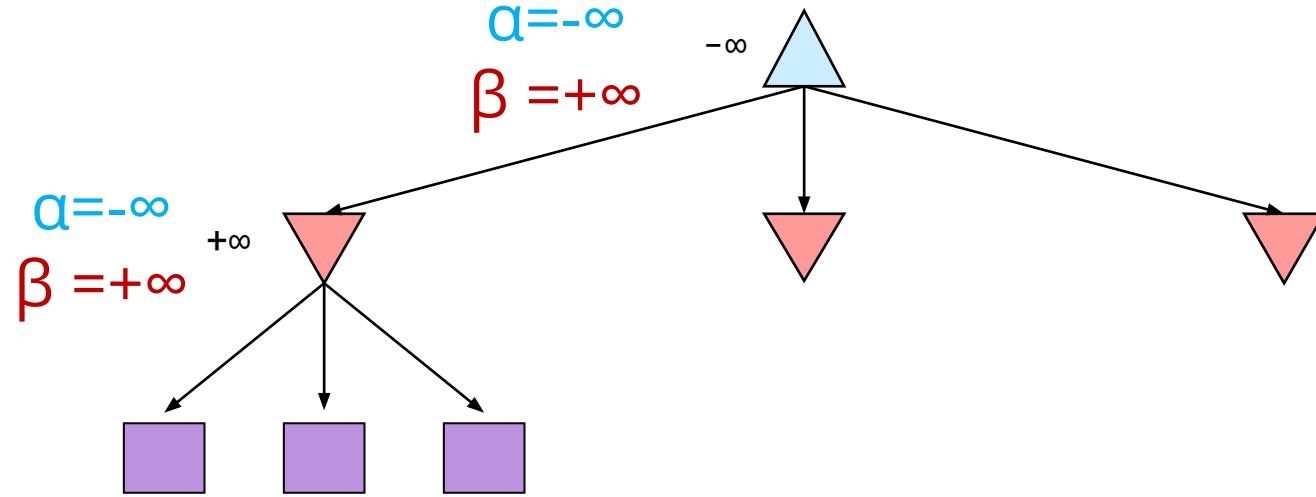
For  $\beta$  think "at most"

# Alpha-Beta Pruning Example

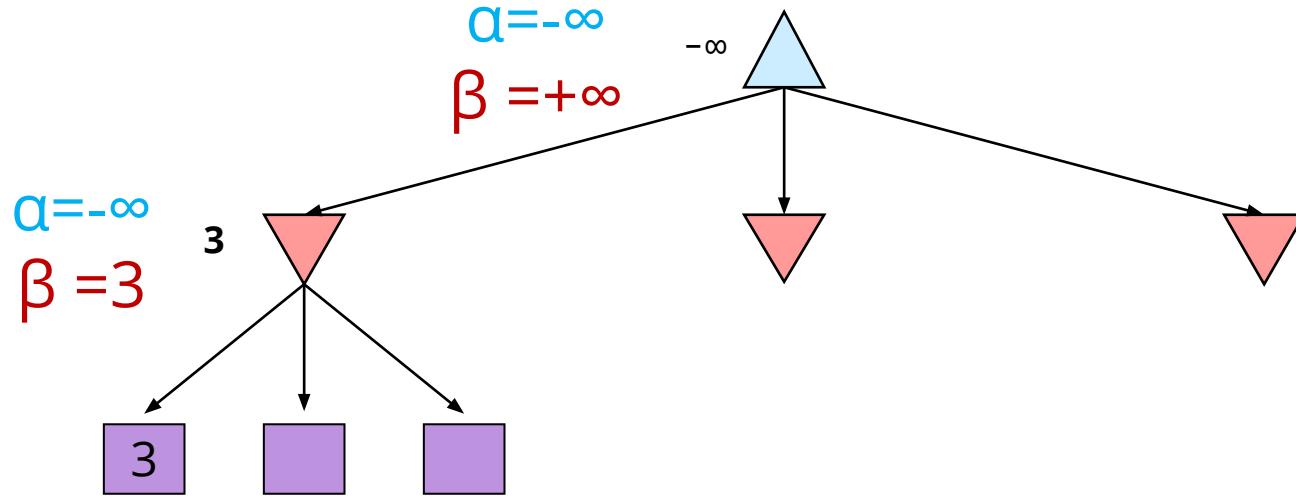
$\alpha = -\infty$   
 $\beta = +\infty$



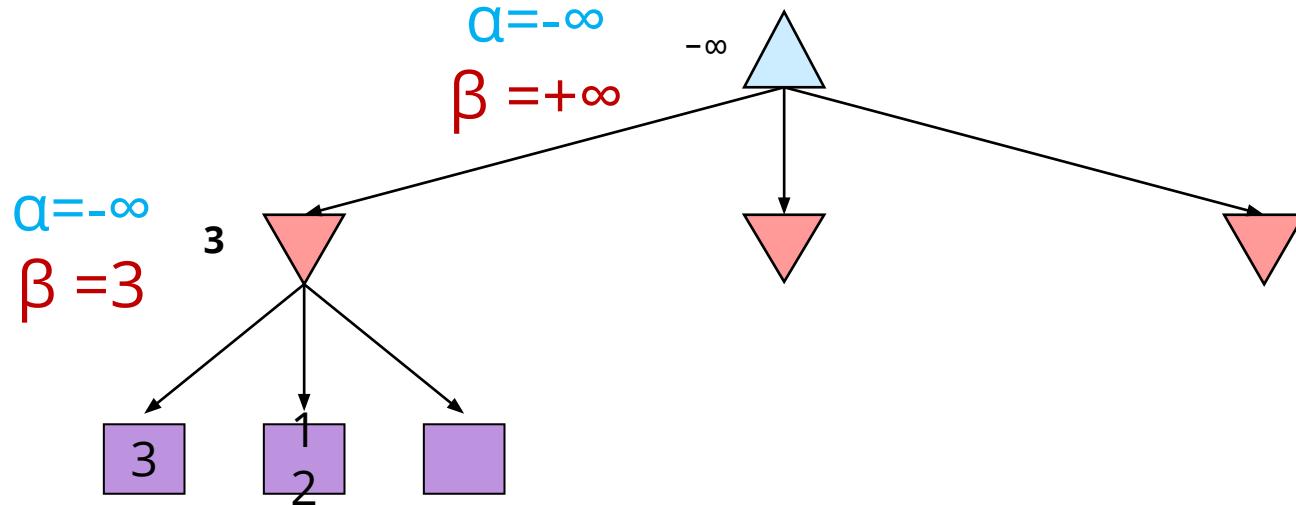
# Alpha-Beta Pruning Example



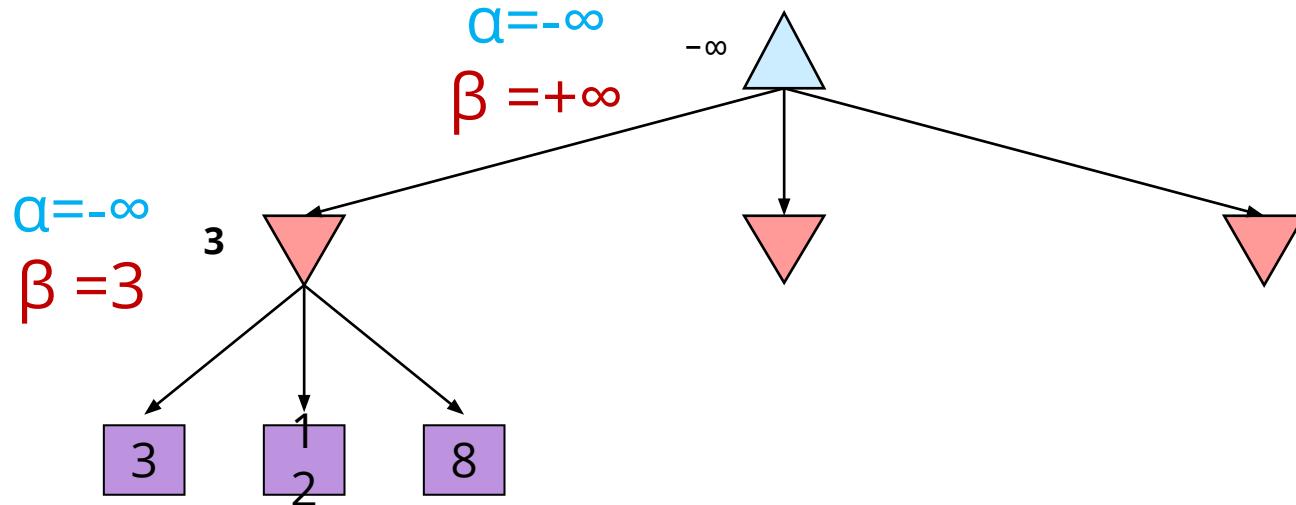
# Alpha-Beta Pruning Example



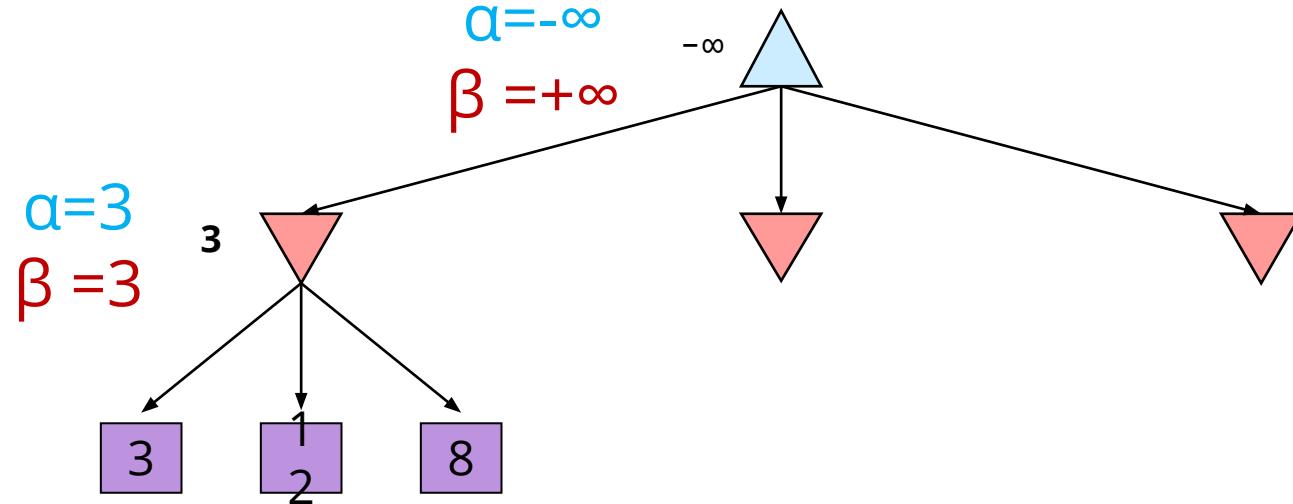
# Alpha-Beta Pruning Example



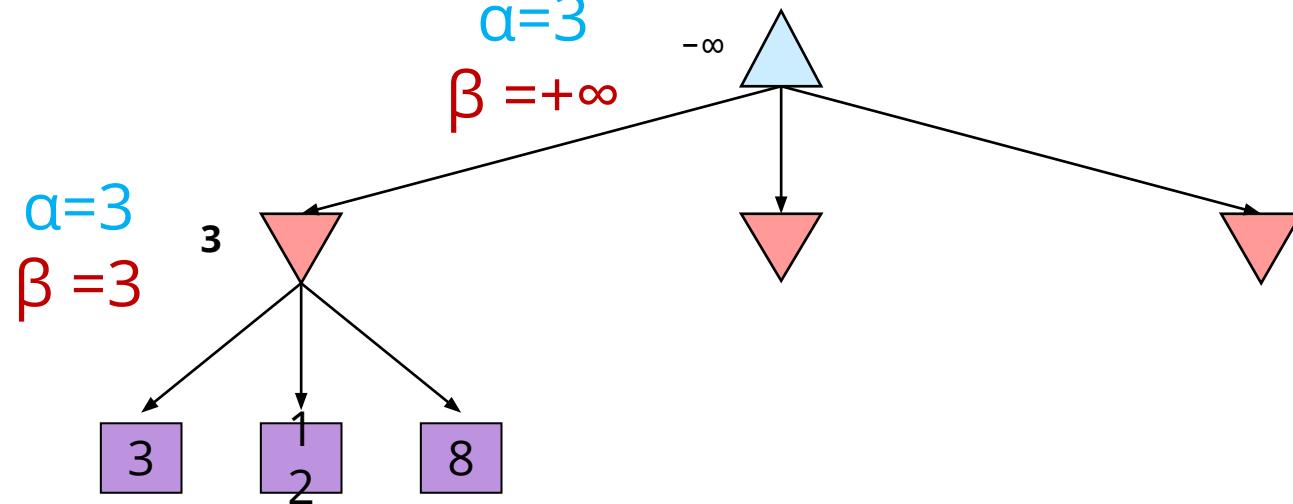
# Alpha-Beta Pruning Example



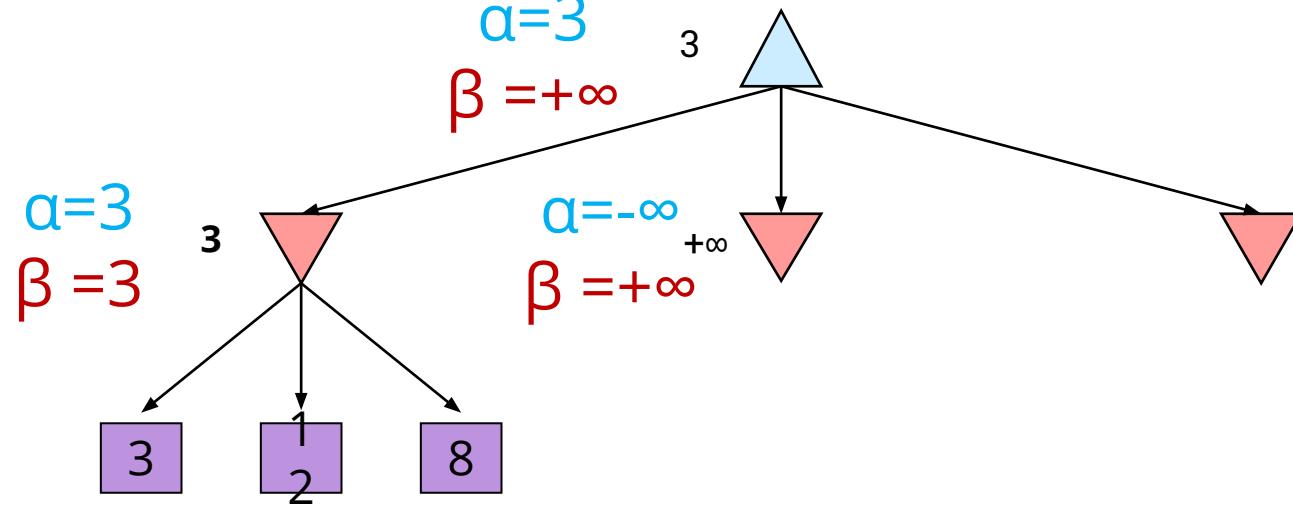
# Alpha-Beta Pruning Example



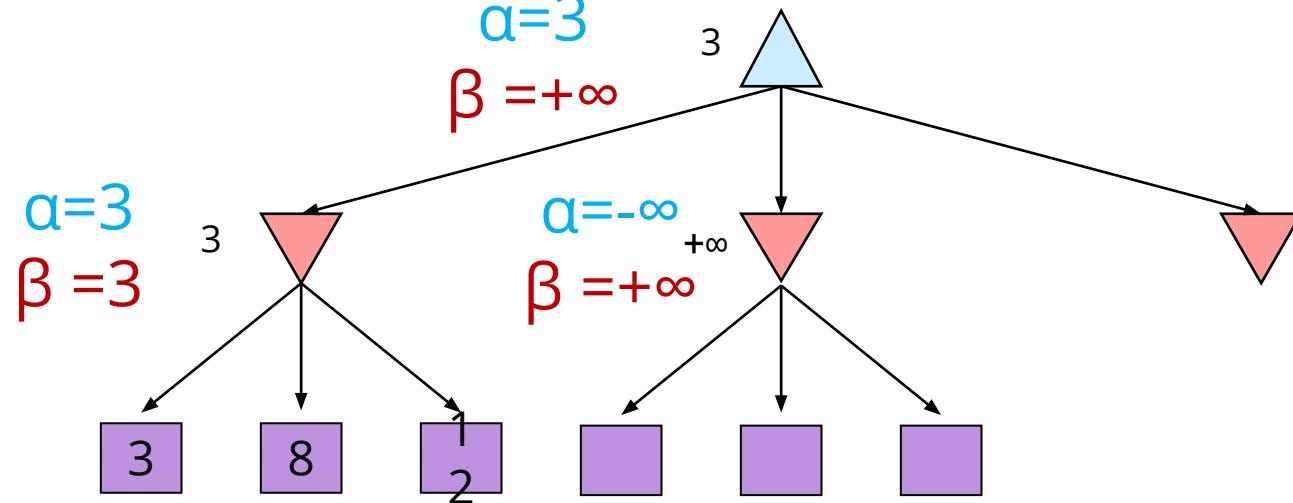
# Alpha-Beta Pruning Example



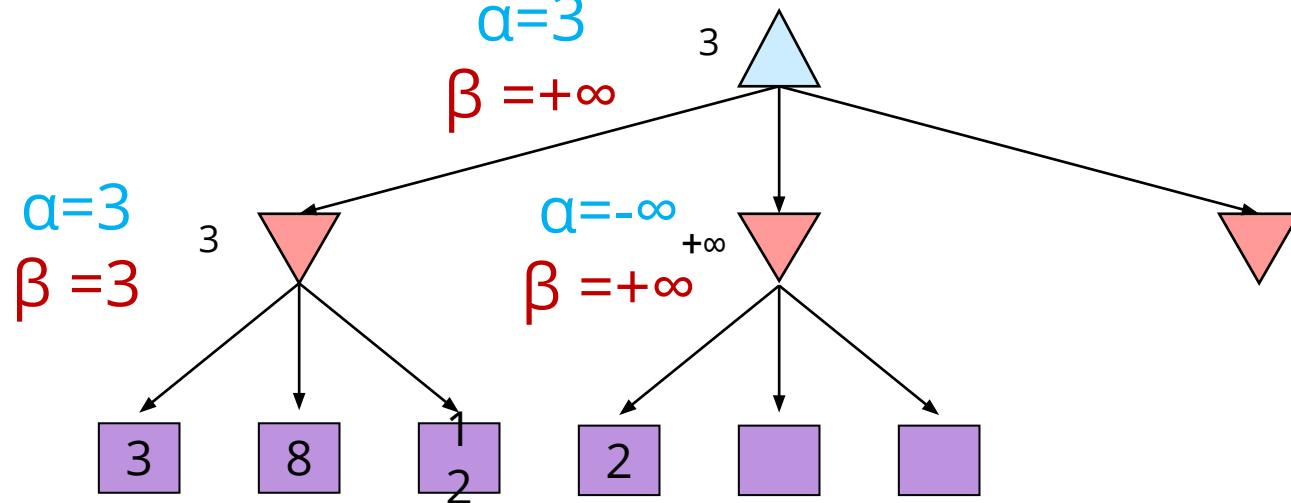
# Alpha-Beta Pruning Example



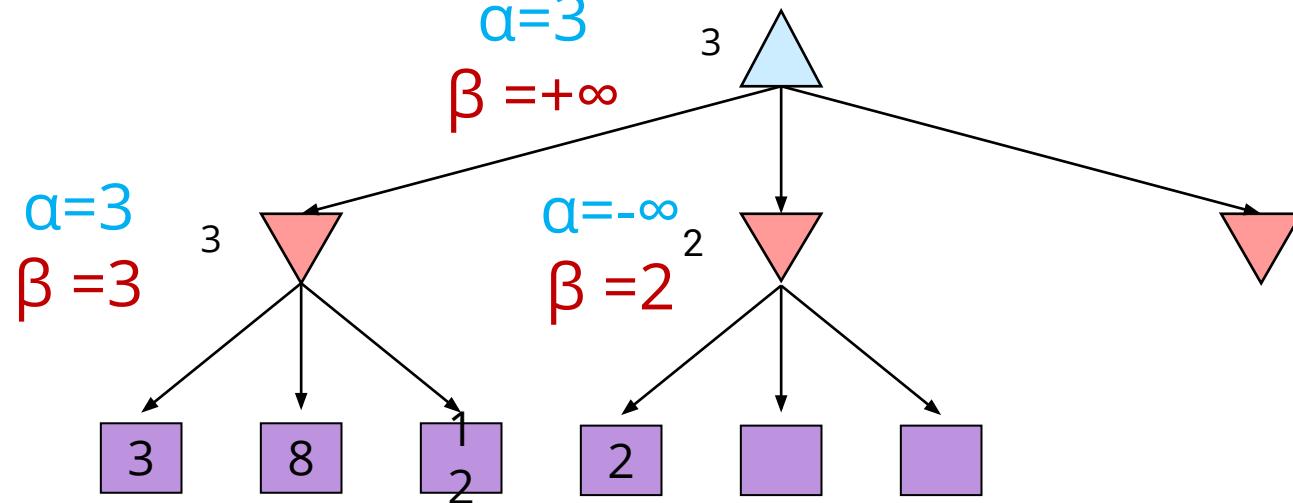
# Alpha-Beta Pruning Example



# Alpha-Beta Pruning Example

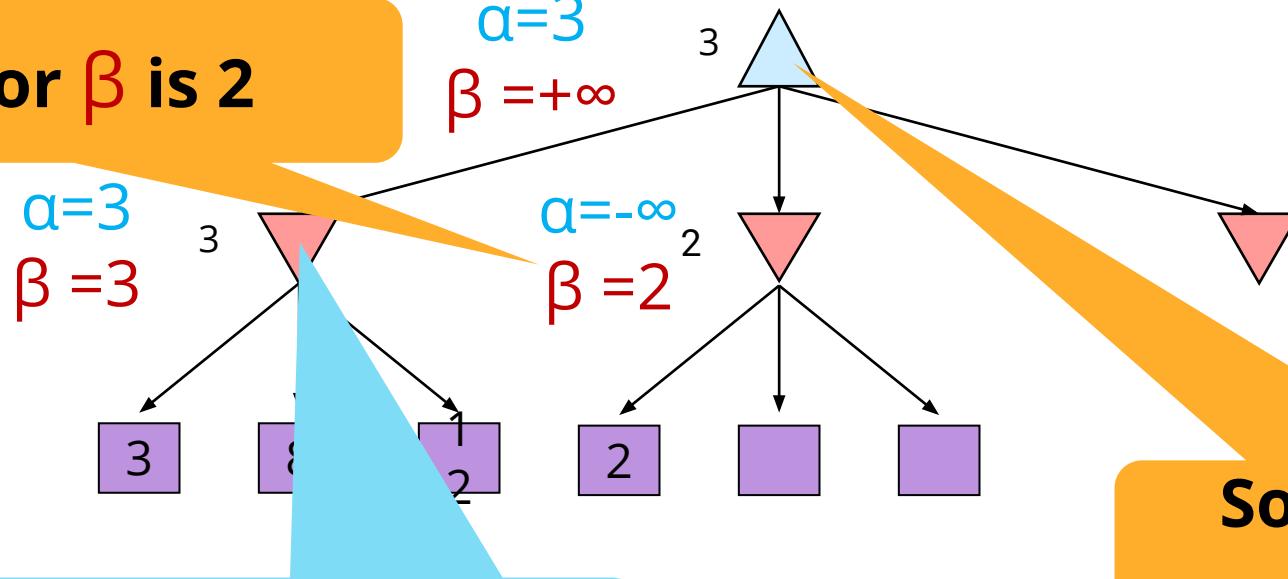


# Alpha-Beta Pruning Example



# Alpha-Beta Pruning Example

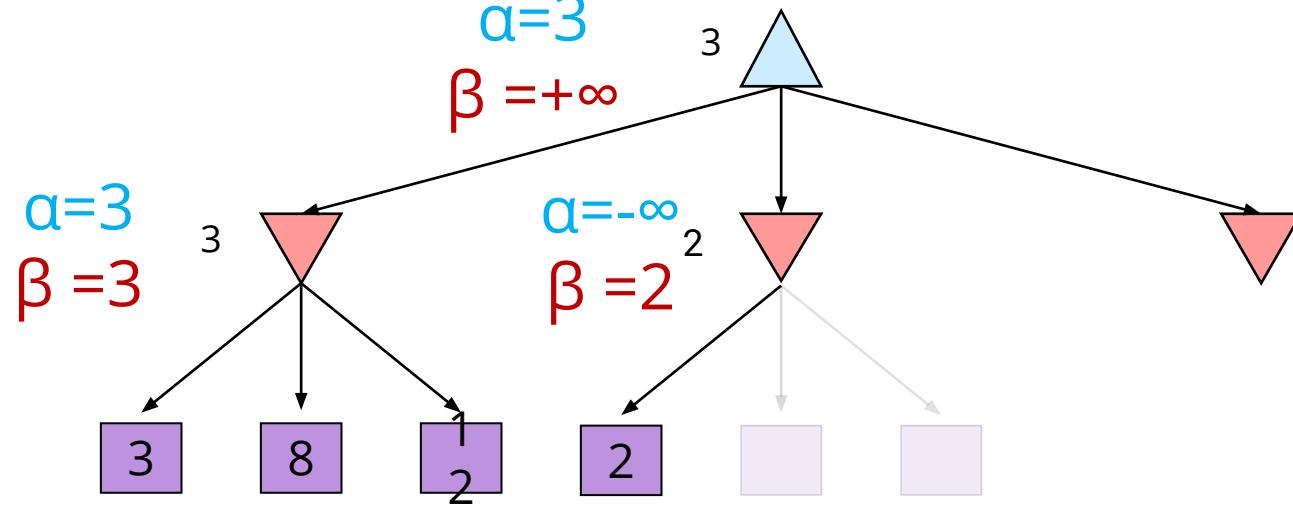
Value for  $\beta$  is 2



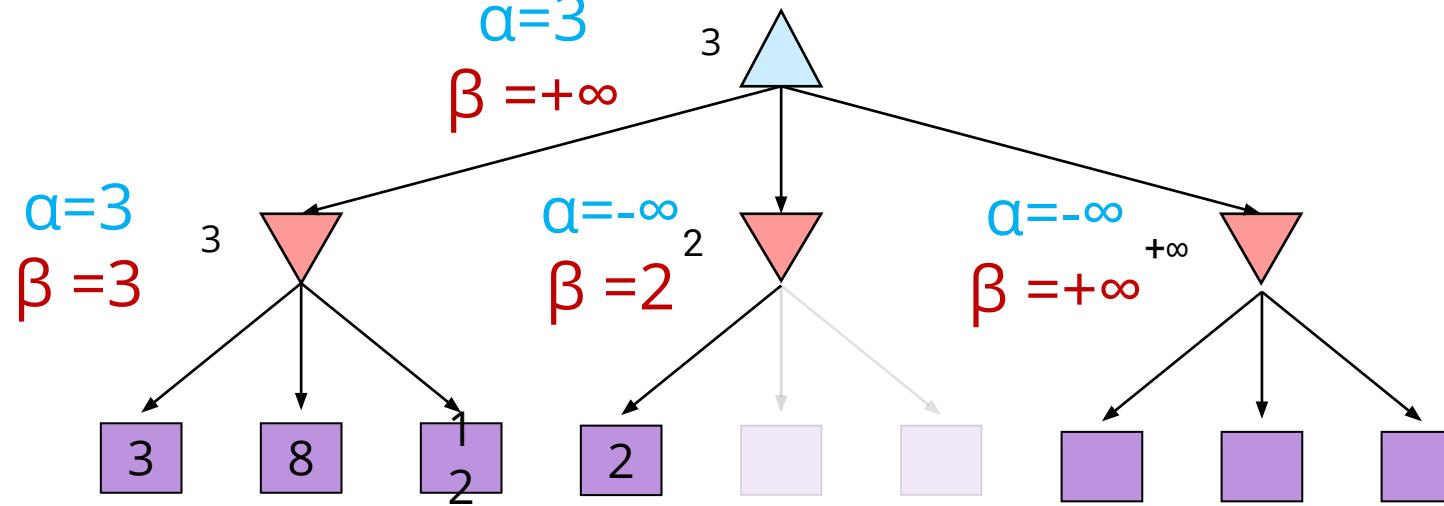
So Max will never choose 2.

But we know that this node is worth at least 3

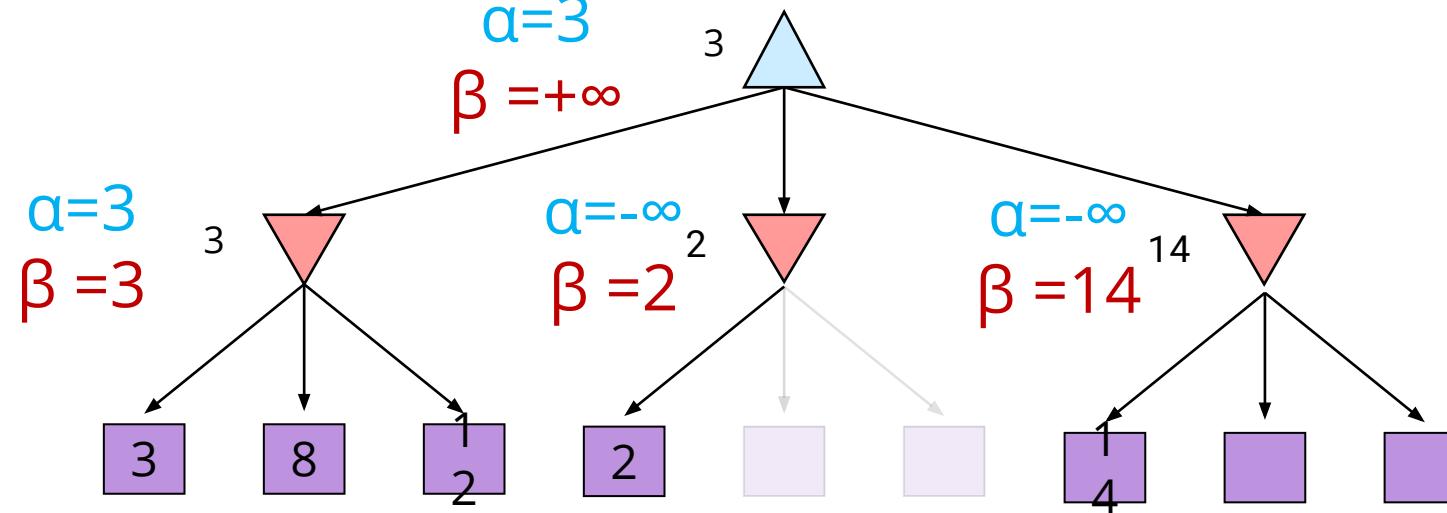
# Alpha-Beta Pruning Example



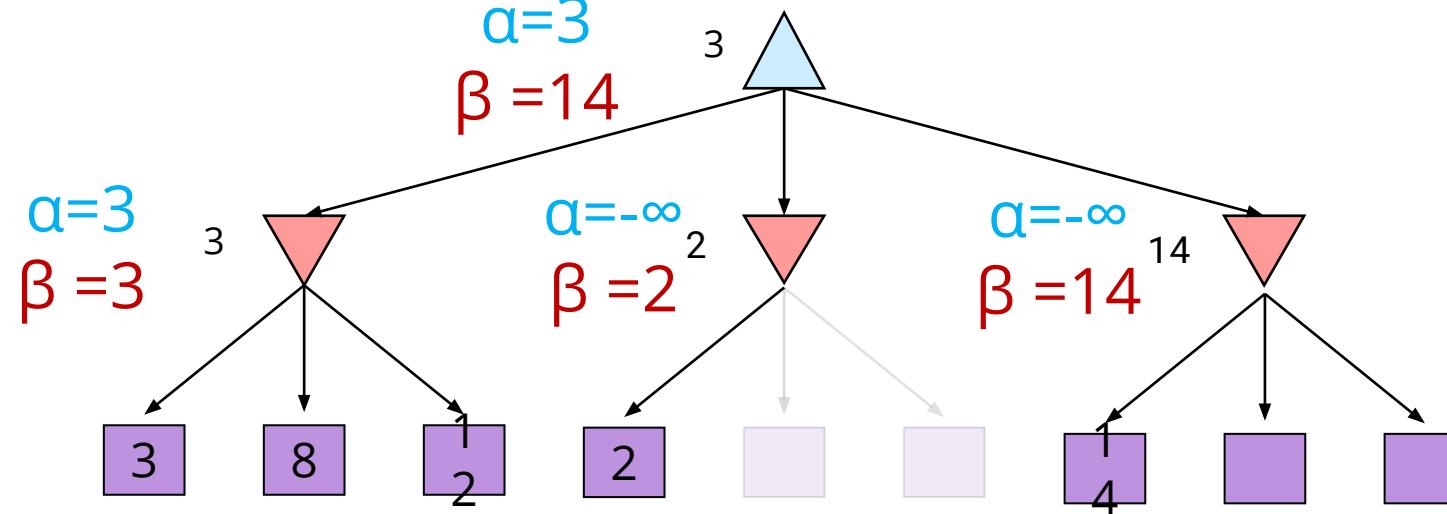
# Alpha-Beta Pruning Example



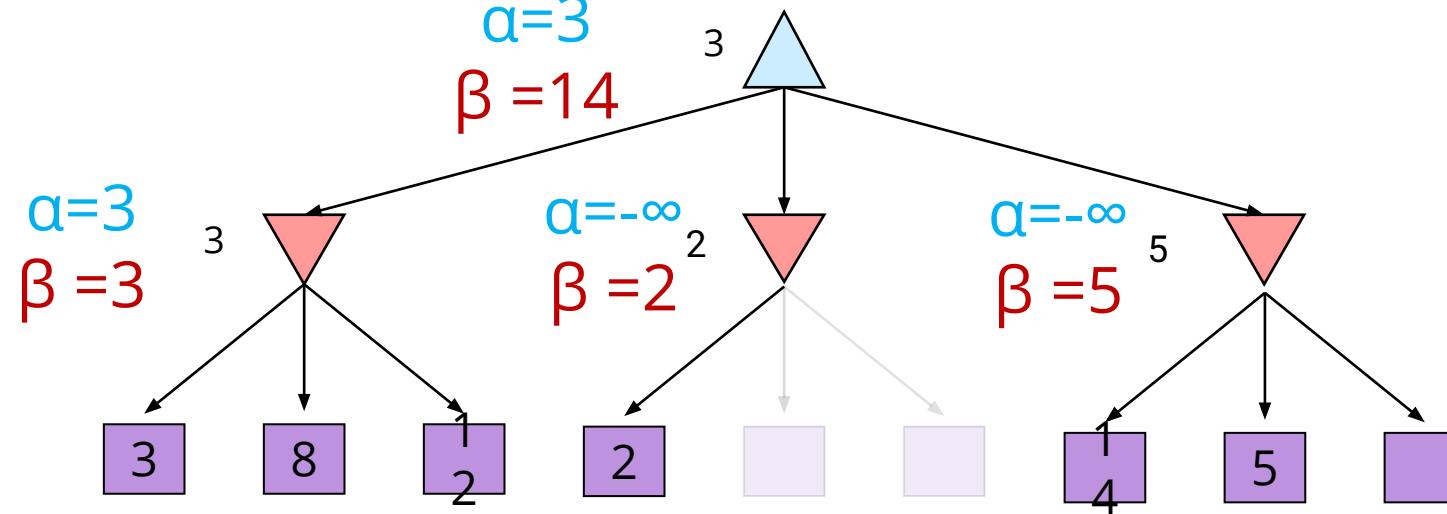
# Alpha-Beta Pruning Example



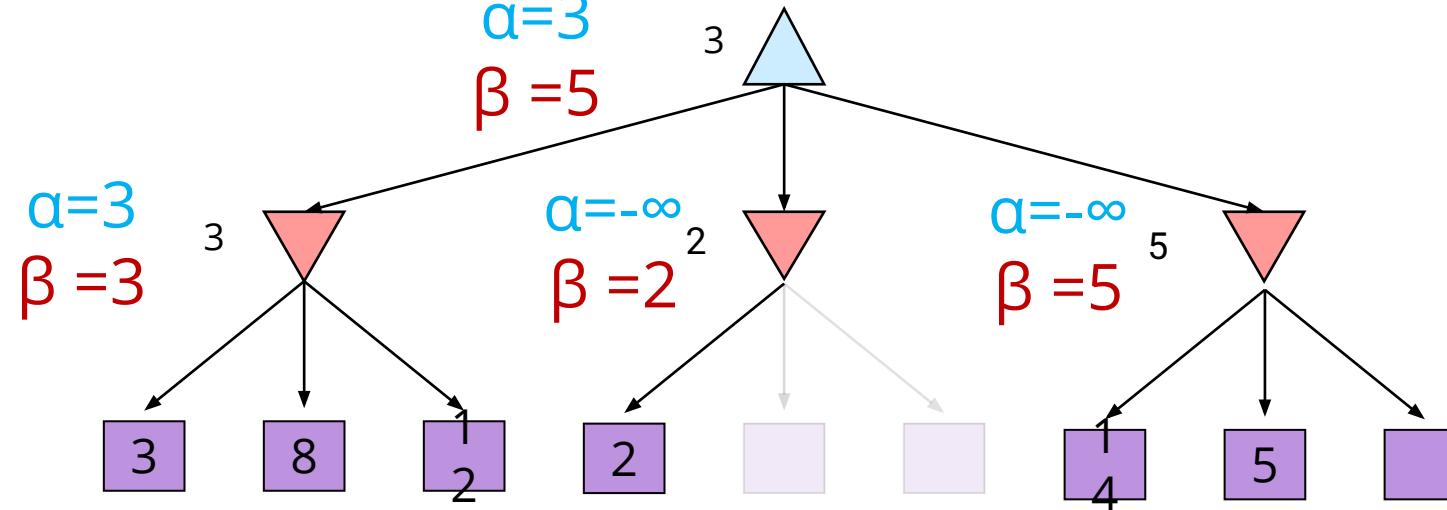
# Alpha-Beta Pruning Example



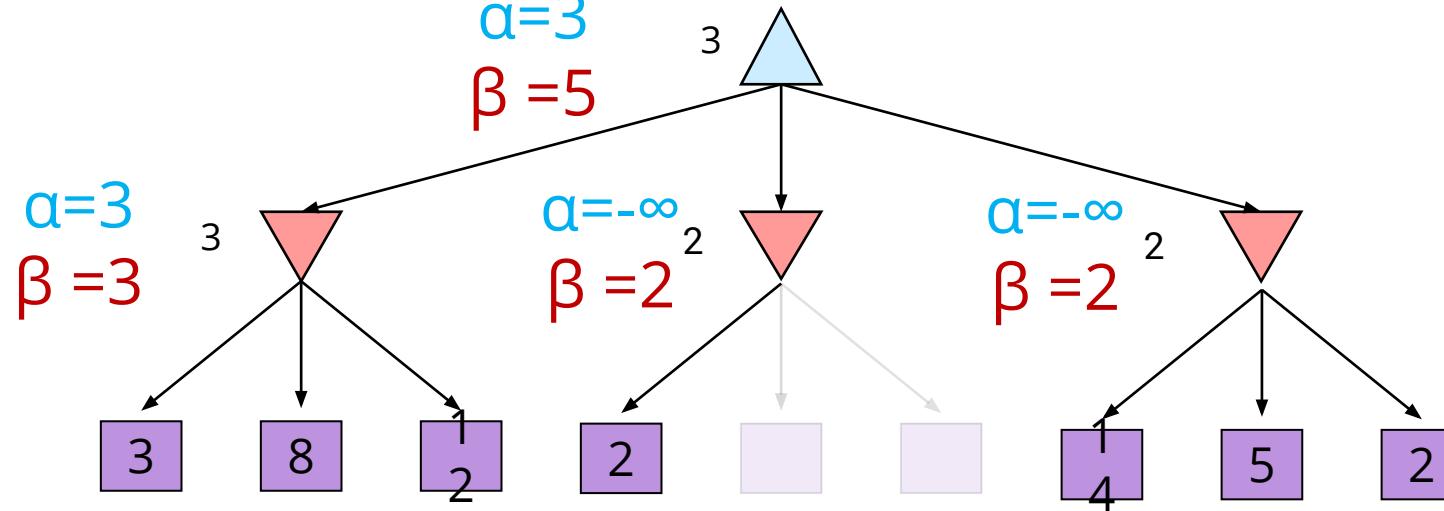
# Alpha-Beta Pruning Example



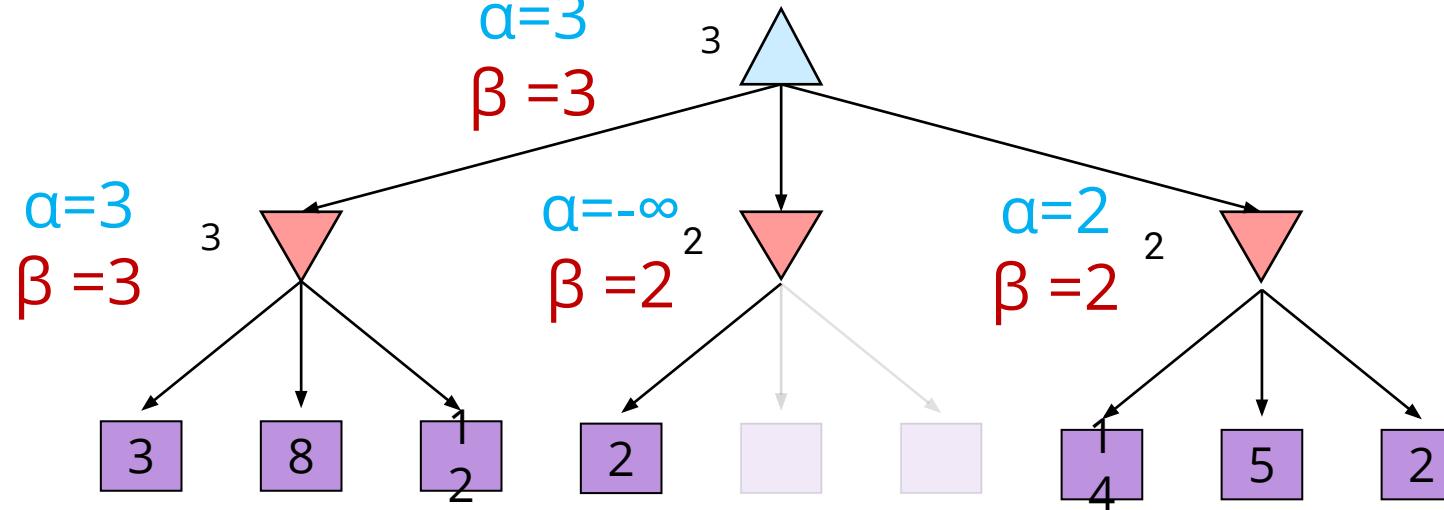
# Alpha-Beta Pruning Example



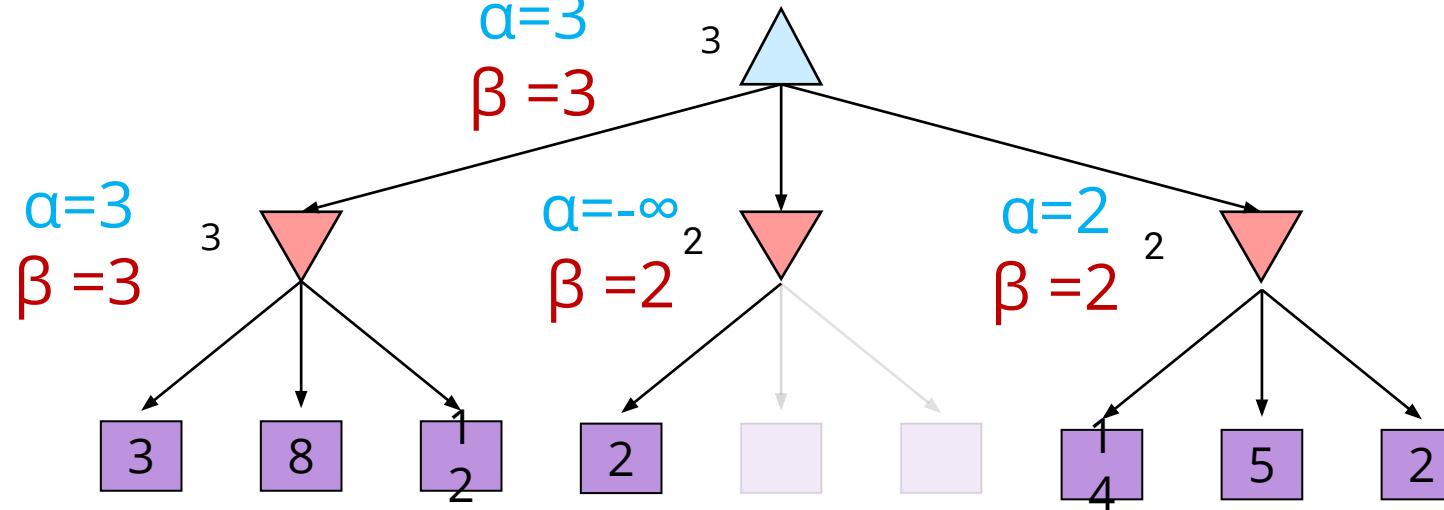
# Alpha-Beta Pruning Example



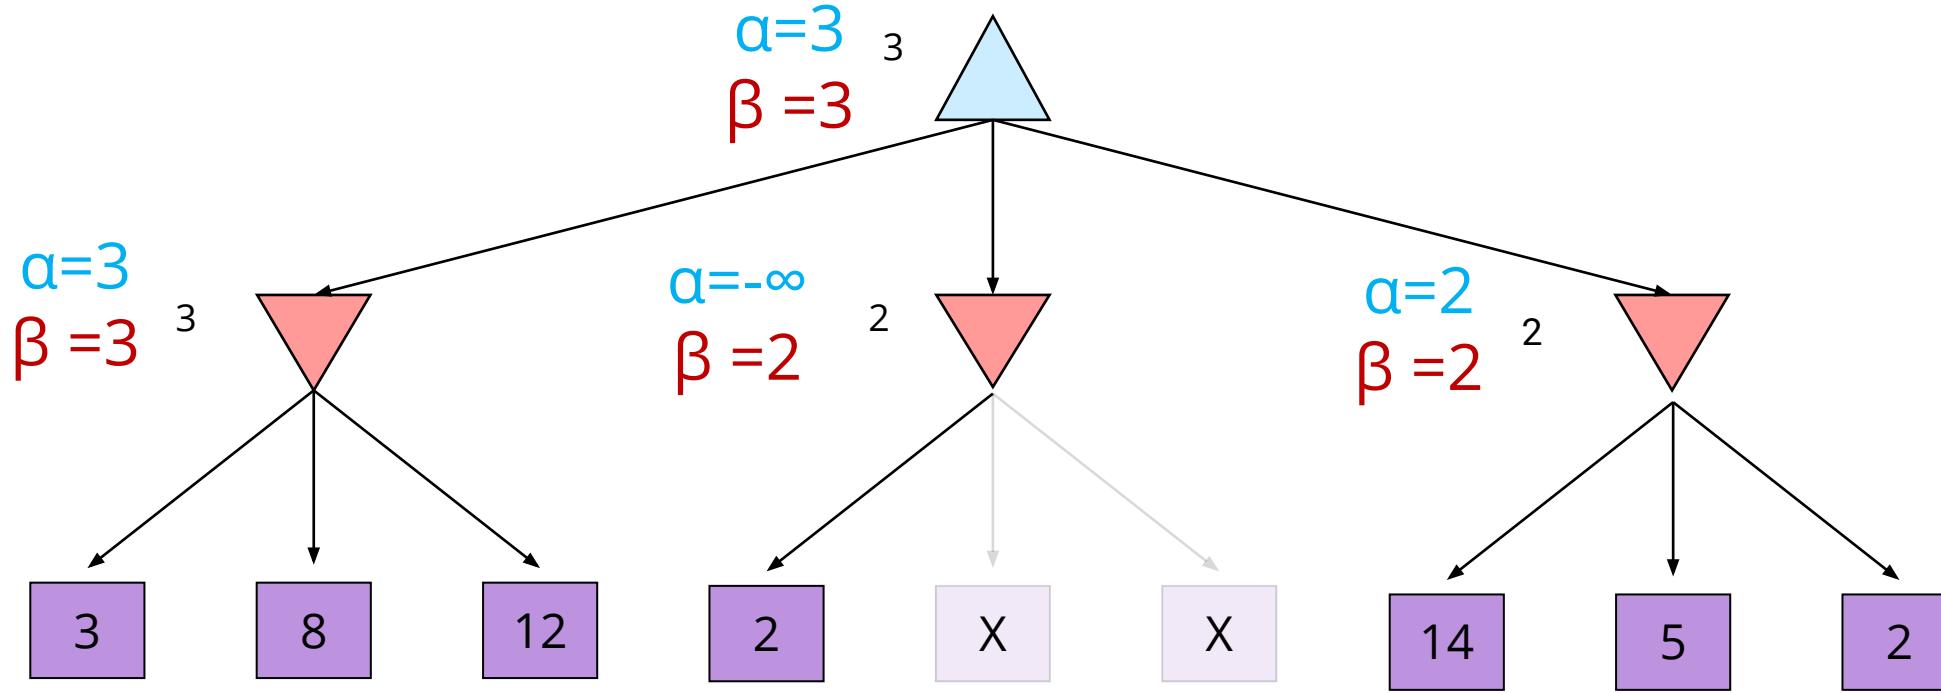
# Alpha-Beta Pruning Example



# Alpha-Beta Pruning Example



# Alpha-Beta Pruning



$\alpha$ : MAX's current lower bound on MAX's outcome  
 $\beta$ : MIN's current upper bound on MIN's outcome

$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

# Review: Evaluation functions

- Evaluates how good a ‘board position’ is
  - Based on *static features* of that board alone
- Zero-sum assumption lets us use one function to describe goodness for both players.
  - $f(n) > 0$  if MAX is winning in position  $n$
  - $f(n) = 0$  if position  $n$  is tied
  - $f(n) < 0$  if MIN is winning in position  $n$
- Build using expert knowledge,
  - Tic-tac-toe:  $f(n) = (\# \text{ of 3 lengths open for MAX}) - (\# \text{ open for MIN})$

(AIMA 5..1)

# Review: Chess Evaluation Functions

- Chess needs an evaluation function since it is impossible to search the game tree deeply enough to reach the terminal nodes
- $f(n) = (\text{sum of } A\text{'s piece values}) - (\text{sum of } B\text{'s piece values})$
- More complex: weighted sum of positional features:
$$\sum w_i \cdot \text{feature}_i(n)$$
- $f(n)$  can be a **weighted linear function**

<b>Pawn</b>	1.0
<b>Knight</b>	3.0
<b>Bishop</b>	3.25
<b>Rook</b>	5.0
<b>Queen</b>	9.0

Pieces values for a simple evaluation function often taught to novice chess players